

Shivam Gupta

shivamgupta@utexas.edu

EDUCATION	University of Texas at Austin	Dec 2024
	Ph.D. in Computer Science Dissertation: Diffusion Models – Training, Sampling, and Reconstruction	
	University of Illinois at Urbana-Champaign	May 2018
	B.S. in Computer Science, Minor in Mathematics	
INTERESTS	Large Language Models, Diffusion Models, Machine Learning, Statistics, related topics	
EXPERIENCE	Google DeepMind	New York, NY
	Research Scientist	Feb 2025 – Present
	• Working on post-training research for Gemini	
	Google DeepMind	San Francisco, CA
	Student Researcher	May 2024 – Oct 2024
	• Worked on leveraging zeroth-order optimization techniques for optimizing data mixtures for training and fine-tuning Large Language Models (LLMs) using evaluation feedback.	
	• Built optimized data and evaluation pipeline from scratch for the Gemma series of models for experimentation using Python, JAX and TensorFlow. Reduced training and evaluation times from several hours for the reference implementation to ≤ 5 minutes using 100+ TPUs. Showed results for algorithmically computing data mixtures to optimize evaluation metrics.	
	• Developed problem-specific optimization techniques and implemented them	
	• Proposed approach for fine-grained sampling of training datasets to optimize evaluation metrics.	
	Massachusetts Institute of Technology	Cambridge, MA
	Visiting Student	Jun 2023 – Aug 2023
	• Proved sharper rates for high-probability mean estimation. Resulted in paper at COLT 2024.	
	University of California, Berkeley	Berkeley, CA
	Visiting Student Researcher	Aug 2022 – May 2023
	• Developed theory for diffusion models, and wrote experiments to improve understanding, resulting in a paper at NeurIPS 2024.	
	• Developed mean estimation algorithms, resulting in multiple papers at NeurIPS, ICML and COLT.	
	Sigma Computing, Inc.	San Francisco, CA
	Research Intern	May 2022 – Aug 2022
	• Worked on designing and implementing anomaly detection algorithms for various datasets	
	University of Wisconsin, Madison	Madison, WI
	Research Intern	May 2020 – Aug 2020
	• Studied gradient descent algorithms to robustly estimate the mean of a high-dimensional Gaussian	
	• Wrote experiments for robust sparse estimation in Python. Resulted in NeurIPS 2022 paper.	
	Jane Street	New York, NY
	Software Developer Intern	May 2016 – Aug 2016
	• Wrote server code to employ state machine replication to send and receive data via RPCs	
	• Developed a market data parser in OCaml	
	Bloomberg L.P.	New York, NY
	R&D Intern	May 2015 – Aug 2015
	• Developed real-time system for Bloomberg Terminal to track messages between services in C++	
PAPERS	14. Posterior Sampling by Combining Diffusion Models with Annealed Langevin Dynamics	
	Zhiyang Xun, Shivam Gupta, Eric Price	
	<i>In submission</i>	
	13. Efficient Knowledge Distillation via Curriculum Extraction	
	Shivam Gupta, Sushrut Karmalkar	
	<i>In submission</i>	
	12. Uncovering the Bias in Diffusion-based Posterior Sampling Methods	
	Shivam Gupta, Brett Levac, Jon Tamir, Eric Price	
	<i>Manuscript</i>	

11. **Faster Diffusion Sampling with Randomized Midpoints: Sequential and Parallel**
Shivam Gupta, Linda Cai, Sitan Chen
International Conference on Learning Representations (ICLR) 2025
10. **Improved Sample Complexity Bounds for Diffusion Model Training**
Shivam Gupta, Aditya Parulekar, Eric Price, Zhiyang Xun
Neural Information Processing Systems (NeurIPS) 2024
9. **Diffusion Posterior Sampling is Computationally Intractable**
Shivam Gupta, Ajil Jalal, Aditya Parulekar, Eric Price, Zhiyang Xun
International Conference on Machine Learning (ICML) 2024
8. **Beyond Catoni: Sharper Rates for Heavy-Tailed and Robust Mean Estimation**
Shivam Gupta, Samuel B. Hopkins, Eric Price
Conference on Learning Theory (COLT) 2024
7. **Minimax-Optimal Location Estimation**
Shivam Gupta, Jasper C.H. Lee, Eric Price, Paul Valiant
Neural Information Processing Systems (NeurIPS) 2023
6. **Finite-Sample Symmetric Mean Estimation with Fisher Information Rate**
Shivam Gupta, Jasper C.H. Lee, Eric Price
Conference on Learning Theory (COLT) 2023
5. **High-dimensional Location Estimation via Norm Concentration for Subgamma Vectors**
Shivam Gupta, Jasper C.H. Lee, Eric Price
International Conference on Machine Learning (ICML) 2023
4. **Finite-Sample Maximum Likelihood Estimation of Location**
Shivam Gupta, Jasper C.H. Lee, Eric Price, Paul Valiant
Neural Information Processing Systems (NeurIPS) 2022
3. **Outlier-Robust Sparse Estimation via Non-Convex Optimization**
Yu Cheng, Ilias Diakonikolas, Rong Ge, Shivam Gupta, Daniel Kane, Mahdi Soltanolkotabi
Neural Information Processing Systems (NeurIPS) 2022
2. **Sharp Constants in Uniformity Testing via the Huber Statistic**
Shivam Gupta, Eric Price
Conference on Learning Theory (COLT) 2022
1. **Nash Equilibrium Computation in Resource Allocation Games**
Shivam Gupta, Ruta Mehta
International Conference on Autonomous Agents and Multiagent Systems (AAMAS) 2018

SELECTED AWARDS

- COLT 2024 Travel Award
- C.W. Gear Outstanding Undergraduate Student Award 2018
- Conference Travel Grant 2018 (for travel to AAMAS)
- Horace and Kate King Wu International Undergraduate Scholarship 2018
- Illinois Engineering Achievement Scholarship 2017
- NTT Data, Inc., Scholarship 2015, 2016
- Franz Hohn and J.P. Nash Scholarship 2015 (for research in scientific computing)
- Perfect score in the Indian National Olympiad in Informatics, and selected as one of 26 students in India to attend the International Olympiad in Informatics training camp 2014
- ACM ICPC Mid-Central Regional: Team placed 6th in 2017, 4th in 2016
- Represented India in SEARCC Software Competition 2013, Colombo, Sri Lanka, and placed 3rd
- Placed 2nd in Dropbox Open programming contest 2015 at UIUC
- Won Bloomberg CodeCon Challenge and invited to CppCon 2015 in Bellevue, Washington
- Won 3Red Trading Tech Challenge in 2015 and 2016 and invited to Chicago

OTHER EXPERIENCE

Research with Prof. Andreas Klockner (UIUC)

Aug 2014 – May 2017

- Developed and implemented mesh-refinement algorithms (in Python and using NumPy) to iteratively refine and coarsen meshes while preserving connectivity information.

	Eventifier - Software Engineer Intern May 2014–Jul 2014 • Developed system to stream Twitter data related to particular topics and classify them using NLP.
	Freelance Programming May 2014–Aug 2014 • Developed Android app for scholarship database company
	Game Development Jan 2013-Aug 2014 • Developed a game engine in C++ and SDL with entity management and rendering functionality • Developed several game prototypes using the engine • Initiated organization of Global Game Jam in 2013 for the first time in India (still occurring annually as of 2025)
SKILLS	• <i>Languages:</i> Python, C/C++, Java, OCaml, Haskell, JavaScript • <i>Software and Libraries:</i> PyTorch, JAX, TensorFlow, NumPy, SciPy, Mathematica, L ^A T _E X
SERVICE	• <i>Program Committee:</i> ALT 2025, 2026; AAAI 2025, 2026 • <i>Reviewing:</i> SODA 2021, 2024; NeurIPS 2023, 2024, 2025; ITCS 2024; ALT 2024; ICLR 2024, 2025; ICML 2024, 2025; TF2M@ICML 2024; AISTATS 2025, 2026; FOCS 2025; MOSS@ICML 2025;
TALKS	• Efficient Knowledge Distillation via Curriculum Extraction April 2025 Google DeepMind • Diffusion Models: Sampling, Training and Reconstruction November 2024 Boson AI • Theoretical Insights for Practical Generative Modeling November 2024 Google Research • Diffusion Models: Sampling, Training and Reconstruction October 2024 Stability AI • Beyond Catoni: Sharper Rates for Heavy-Tailed and Robust Mean Estimation July 2024 Conference on Learning Theory 2024 (Online) • Sample-Efficient Training for Diffusion November 2023 IFML Workshop on Generative AI, UT Austin • A Finite-Sample Theory for Mean Estimation with Fisher Information Rate October 2023 MIT Algorithms and Complexity Seminar • A Finite-Sample Theory for Mean Estimation with Fisher Information Rate October 2023 CMU Theory Lunch • Finite-Sample Symmetric Mean Estimation with Fisher Information Rate July 2023 Conference on Learning Theory 2023, Bangalore, India • High-dimensional Estimation via Norm Concentration for Subgamma Vectors July 2023 International Conference on Machine Learning 2023 (Online) • Sharp Constants in Uniformity Testing via the Huber Statistic July 2022 ML Tea, UT Austin
TEACHING	At UT Austin: • <i>Teaching Assistant</i> , Sublinear Algorithms (CS 395T) Fall 2020 • <i>Teaching Assistant</i> , Introduction to Algorithms (CS 331) Spring 2019 • <i>Teaching Assistant</i> , Machine Learning (CS 395T) Summer 2019, Fall 2019 • <i>Teaching Assistant</i> , Elements of Data Analytics (CS 329E) Spring 2019 • <i>Teaching Assistant</i> , Elements of Computer Programming (CS 303E) Fall 2018 At UIUC: • <i>Course Assistant</i> , Introduction to Algorithms (CS 374) Spring 2018 • <i>Grader</i> , Algorithms II (CS 473) Spring 2018
RELEVANT COURSES	Probability and Stochastic Processes, Learning Theory, Coding Theory, Theoretical Statistics, Randomized Algorithms, Markov Chains and Mixing Time, Approximation Algorithms, Combinatorial Mathematics, Numerical Linear Algebra, Wireless Networking